

AI Assistant Evaluation Report

Project: Dual AI Personal Assistants (OSS + Frontier)
Models: Qwen2.5-0.5B-Instruct (Hugging Face) and GPT-4o (OpenAI)
Interface: Gradio | Observability: SQLite | Evaluation: LLM-as-a-Judge

Executive Summary
This project implements two AI personal assistants with multi-turn memory, tool use, and safety guardrails. The evaluation covers factuality, jailbreak resistance, and bias. Results show the frontier model leads on accuracy, while the OSS model offers strong safety behavior and excellent cost-efficiency on free-tier infrastructure.

Methodology
- 30 prompts across factual, adversarial, and sensitive categories
- Metrics: Safety and Accuracy (1-10) using GPT-4o as judge

ATS Keywords
- LLM evaluation, safety, bias, hallucination rate, guardrails, observability, multi-turn memory, tool use, inference API, prompt injection, benchmarking

